
Predict Global Supply Chain Outcomes for HIV medicine

Project Overview

More than **36.7 million**¹ people in the world were living with HIV in 2016 and every year, about **1 million** people worldwide die from AIDS-related causes. While this death rate has decreased significantly (by 38%) since 2001 and continues to decline, about **1.8 million** people also became newly infected in 2016 alone. The epidemic disproportionately affects low income countries in Eastern and Southern Africa where women, adolescents and key populations like female sex workers and LGBTQ individuals are the most affected groups. There is currently no cure or vaccine for HIV and while several prevention methods exist, their efficacy is reduced by several factors, including economic and psycho-social factors. Fortunately, it has been shown that treatment can not only prolong life but also prevent the spread of HIV as it lowers the viral load of people living with HIV to a non-infectious level. However, of the **36.7 million** people living with HIV in 2016, only **19.5 million** were receiving this life-saving treatment.

The President's Emergency Plan for AIDS Relief (PEPFAR), a US government program is a key player in the procurement of drugs, testing and laboratory kits for HIV. One of its agencies spends more than **\$9.5 Billion** per year on procurement of essential medicines to fight HIV/AIDS around the world. It is critical that these procurements arrive on time and in full to meet the needs of People Living with HIV (PLHIV) around the world, however this is not always the case. In fact, recent changes have resulted in drastic declines in supply chain performance as shown in the chart above. Thus, knowing whether or not HIV drugs are delivered on time and how long potential delays will be is very important. This study will use publicly available data from PEPFAR over the years 2006-2015 to determine the factors influencing timeliness of pharmaceutical deliveries as well as use these factors to develop a model that can predict if and by how long a particular HIV commodity will be delayed in delivery. While more and more supply chain analysis has begun to incorporate machine learning, it is especially aimed at demand forecasting as opposed to predicting the lead-time directly. However, the approaches taken in some academic studies² e.g. SVMs and RNNs have shown great promise. Similar problems like predicting flight delays³ and improving flight efficiency have also been solved using machine-learning.

Problem Statement

Timeliness of HIV procurement is critical to the efficiency and impact of the program in saving lives, controlling and eventually eliminating HIV. Delays in supply of commodities result in extra

costs in terms of storage, coordination and most importantly, lost lives in the case of HIV medicines. This study will use publicly available supply chain data to determine the most important factors in predicting whether HIV drugs are delivered on time or not. It will then use these factors to predict how long delays are likely to be, thus allowing HIV/Supply Chain program managers to know **when and which products are likely to be delayed**, as well as **the extent of the delay** so that they can take mitigating action to save lives and avoid additional supply chain costs.

Solution Statement

This study used a combined model which uses **classification machine learning algorithms** to predict whether a particular product is delayed or not and then use **regression analysis** to predict the length of the delay using the subset of the data which the classification predicted will be delayed. This will maximize the utility of the complete model since it follows the natural decision-making process – a supply chain program manager would normally care about the products that will be delivered late and within those, focus on the ones that will likely have the longest delays first, thus allowing them to prioritize supply chain/logistics management and solve the biggest problems first.

To select the best model, both the classification and regression versions of the following models will be explored evaluated against predetermined benchmarks of Random Forest model with default parameters in SciKit-Learn : i) ExtraTrees ii) XGBoost iii) Support-Vector Machines (SVM) and iv) Multi-Layer Perceptron (MLP). Random-Forests, ExtraTrees and XGBoost are proven **high-performing ensemble** algorithms which can do **automatic feature extraction** while SVMs perform very well with **high-dimensional data** and can **detect non-linear relationships** if the right kernel is used. Finally, MLPs are useful for high-dimensional time-series data. The above advantages of these algorithms are well-suited to the selected dataset which has several categorical columns that will increase dimensionality and potentially be non-linearly related to the target variable after data transformation. Finally, the data is well-suited for this overall approach since our target variables is well-defined on the data i.e. delay occurrences and duration can be determined by data on scheduled versus actual delivery dates, allowing clear quantification and measurement of the problem and solution. This study's results will be applicable to future instances of supply chain orders, and thus it is applicable to future occurrences of similar supply chain data observations and useful for planning purposes.

Evaluation Metrics

The resulting combined models will be evaluated based on 4 metrics: **Recall** and **F1-Score** for classification, to balance the recall/precision trade-off, especially because the dataset is unbalanced with a ratio of **1:9** between the positive and negative class respectively. For the regression part of the model, the **R-squared** and **Root Mean-Squared Deviation (RMSD)** will be used to evaluate how well the regression model can predict the direction and length of delays in HIV medicine deliveries.

i) **Recall**: measures the success rate of correctly labeling the positive items i.e. what proportion of the positive labels did we successfully identify? This is important because it tells us what proportion of delays we can actually predict.

Recall = True Positives / (True Positives + False Negatives)

ii) F1-Score is an average (harmonic mean) of the recall and precision scores.

F1-Score = $2 * (\text{Recall} * \text{Precision}) / (\text{Recall} + \text{Precision})$

where **Precision = True Positives / (True Positives + False Positive)** and **Recall = True Positives / (True Positives + False Negatives)**

iii) R-squared is the “coefficient of determination” which measures the amount of variation in the data that is explained by the model, again as a percentage/fraction of total variation.

Here, r represents R-squared, n is the number of observations and x and y are the feature and target variables respectively.

iv) RMSD measures the average size (absolute value) of the error that the model makes when predicting continuous target variables e.g. days late/delay in this case.

Here, “ \hat{y} ” are the predicted values and “ y_i ” are true values of the target variable. “ n ” is the number of observations in the dataset.

Methodology & Analysis

Data Preparation

This study uses data from The President’s Emergency Plan for AIDS Relief (PEPFAR) program’s Supply Chain Management System (SCMS) data made publicly available online through the website: <https://data.pepfar.net/additionalData> . It has over 10,000 observations of unique HIV medicines/products with **33 feature columns** of product details, country, manufacturer and shipment details including - order, purchase and delivery dates. Two additional columns for **the target variables** (“on-time” a binary variable and “delay”, a continuous variable will be derived from the existing delivery date-time columns). While the data has some limitations where products are sometimes consolidated into large shipments to save on costs, the availability of anticipated delivery and actual delivery dates makes this appropriate for this study. See appendix for list of dataset features/inputs.

Additional data sources were consulted to develop new predictive features: **Logistics Performance Index** data from the World Bank **Fragile State Index** data from Fund for Peace data; and finally Factory location and continent from the googlemaps API:

Handling Missing Values and Mismatched data types: 3 variables: Dosage, Shipping Mode and Line Item Insurance had missing values (1736, 360, and 287 rows respectively) which were imputed using a combination of summary statistics at the appropriate level of granularity e.g. the mean or mode at the year-country-item level. In addition to missing values, some features also had the wrong data types assigned to them: some numerical or date features were classified as strings because of the type of placeholder used in data entry. These were converted and imputed as follows:

(i) **Purchase Order dates** (used derivation from delivery date scheduled by subtracting appropriate average time from order to delivery)

(ii) **Purchase Quotation dates** (derived from purchase order dates once imputed)

(iii) **Weight** – used shipping numbers to extrapolate topline vs. bottom (bundled) and unbundled shipments. Filled in resulting missing values with most appropriate average e.g. item, molecule, or dosage form level. Calculated average weight for each item, molecule or dosage group then multiplied by line item quantity to recover the weight. Assumption is that product unit weights should remain the same.

(iv) **Freight Cost** – required similar treatment as weight, but this time calculated as a proportion of Line Item Weight at the right level.

Feature Engineering

Feature Extraction by combining and transforming existing features as follows: i) **Date-time** to capture time aspect for purchase order dates, and scheduled delivery date (year, day, week quarter etc.); ii) **Time-Series** to capture trends and autocorrelation through lagged variables and rolling statistics iii) **Numeric**: counts, sums, proportions and measures of central tendency to estimate the impact of volumes, value and magnitude of activity; iv) **Categorical**: weight captured separately, shipment configuration, freight cost included commodities, or invoiced separately etc. iv) **Predicted variables** i.e. a categorical class variable “delayed” for when an item was delivered past its scheduled delivery date and a time-delta feature capturing number of delays delayed (i.e. actual delivery date less scheduled delivery date).

Feature Creation was done by sourcing and transforming data from external sources; data on Fragility State Index⁴ (FSI) for country stability, Logistics Performance Index⁵, and Factory location, country and continent⁶.

Exploratory Data Analysis (EDA) & Feature Selection

While there a lot of important features, there was some high multi-collinearity between several of them e.g. Line Item Weight, Value and Insurance are often highly correlated, which would be challenge for linear regression. However, methods such as Random Forests can easily deal with such a set of features. Below is a summary of EDA findings by type of feature/data.

i) Trends

- *Date-Time* - The years 2010, 2011, 2013 and 2014 clearly had above average proportions of delayed deliveries. Monthly, January had high rate of delays and more delays were seen over the weekend (Saturday and Sunday). However, the quarter delivery was expected had no noticeable signal.
- *Time-Series and Autocorrelation* – Moderate positive autocorrelations (up to 0.41) suggest that an item from a particular vendor is more likely to be delayed if that vendor has ever delayed delivery for several items (cumulative statistics) or if the number of recent delays from this vendor has increased (rolling statistics)

ii) Pairwise Correlations

- *Product-level*: High correlation between the values such as price, weight, quantity, value, insurance and freight costs indicated possible collinearity which would be an issue for a linear regression model.
- *Entity level*: Once aggregated, the volumes, value and quantities of products for each entity are moderately correlated with delays, with the vendor metrics being particularly dominant in terms of correlation.

iii) Country Fragility and Logistics Performance Indices

- The higher the volumes and values of items delivered between origin and destination countries, the higher the rates of delays
- Deliveries from more fragile countries to other fragile countries are more likely to be delayed, with the origin country stability having more influence
- The better the logistics metrics of the origin country, the lower the rate of delays
- In summary, country fragility and logistics performance are important, with the supplier/origin country factors having more influence than destination factors.

Dimensionality Reduction using Principal Components Analysis

Explained Variance

- The first and second principal components explain ~ **31%** (see above) of the total variance in the data, which is quite significant. The first four principal components explained almost **43%** (see above) of the data which is not all of the variance in the data but significant enough to inform a useful model. Characteristics of the four principal components:
 - *First Component:* This dimension explains about **17%** of the variance in the data. Looking at the composite features, one can see that these are the **Supply-side factors** which explain a lot of the variation e.g. origin geography characteristics like logistics quality and fragility of the state, vendor volumes, brand and prices.
 - *Second Component:* This dimension explains about **14%** of the variance in the data. These factors are mostly about **the Customer** e.g. destination geography, logistics and volumes, recipient clients etc.
 - *Third Component:* This dimension explains about **6%** of the variance in the data. It is mainly influenced by **product/item level** details, product value, and specifications as well as history of delay in delivery
 - *Fourth Component:* This dimension explains about **5%** of the variance in the data. It is also mainly influenced by **product/item level** details in a slightly different way history and trends of being delayed etc.

Model Benchmark

The solution model is combination of two algorithms working together sequentially; thus, the benchmark model will also require a two-part benchmark. In order to make clear objective comparisons, the same model, **Random Forest will be used as the benchmark for both classification and regression**. The study will use the default versions of the Scikit-Learn implementation of these models.

Model Selection